

1 Graph Search And Connectivity

1.1 A Few Motivations

1. Check if a network is connected (can get to anywhere from anywhere else)
2. driving directions
3. formulate a plan [e.g. how to fill in a sudoku puzzle]
 - (a) nodes = a partially completed puzzle
 - (b) arcs = filling in one new square
4. compute the “pieces (or components)” of a graph
 - (a) clustering, structure of the web graph, etc.

Goals:

1. find everything findable from the a given start vertex
2. don't explore anything twice $O(m + n)$ time

Generic Algorithm (given graph G , vertex s)

- initially s explored, all other vertices unexplored
- while possible:
 - choose an edge (u, v) with u explored and v unexplored (if none, halt)
 - mark v explored

Claim: at end of the algorithm, v explored \iff has a path from s to v .

Proof: induction on number of iterations

by contradiction

Suppose G has a path P from s to v :

but v is unexplored at the end of the algorithm. Then \exists edge $(u, w) \in P$ with u explored and w unexplored. But then algorithm would not terminate, contradiction.

1.1.1 BFS vs. DFS

Note: how to choose among the possibility many “frontier” edges?

Breadth First Search (BFS)

$O(m + n)$ using a queue (FIFO)

explore nodes in “layers”

can compute shortest paths

can compute connected components of an undirected graph

Depth First Search (DFS)

$O(m + n)$ time using a stack (LIFO), or via recursion

explore aggressively like a maze, backtrack only when necessary

compute topological ordering of directed acyclic graph

compute connected components in graphs

1.2 Breadth-First Search (BFS): The Basics

1.2.1 The Code

```
BFS(graph G, start vertex s)
[all nodes initially unexplored]
- mark s as explored
- let Q = queue (FIFO), initialized with s
- while Q != 0
    - remove the first node of Q, call it v
    - for each edge (v, w):
        if w unexplored
            - mark w as explored
            - add w to Q (at the end)
```

Claim #1: At the end of BFS, v explored \iff G has a path from s to v .

Reason: Special case of the generic algorithm

Claim #2: running time of main while loop = $O(n_s + m_s)$, where:

- n_s = number of nodes reachable from s
- m_s = number of edges reachable from s

1.2.2 BFS and Shortest Paths

Goal: Compute $dist(v)$, the fewest number of edges on a path from s to v .

Extra code:

initialize $dist(v) = \begin{cases} 0 & \text{if } v=s \\ +\infty & \text{if } v \neq s \end{cases}$

when considering edge (v, w) :
- if w unexplored
- set $dist(w) = dist(v) + 1$

Claim: at termination, $dist(v) = i \iff v$ in i^{th} layer (i.e. \iff shortest $s - v$ path has i edges)

Proof idea: every layer i node w is added to Q by a layer $(i - 1)$ node v via the edge (u, w)

1.2.3 BFS and Undirected Connectivity

Let $G = (V, E)$ be an undirected graph.

Connected components = the “pieces” of G .

Formal definition: equivalence classes of the relation $u \sim v \iff \exists u - v$ path in G .

Equivalence relation:

satisfy these three properties

- reflexive - everything needs to be related to itself
- symmetric - if $u \sim v$ then $v \sim u$, true since undirected graph
- transitive - if $u \sim v$ and $v \sim w$ then $u \sim w$

Goal: compute all connected components

Why: check if network is disconnected

- graph visualization - clustering

1.2.4 Connected Components via BFS

To compute all components (undirected case)

```

- all nodes unexplored [O(n)]
[assume labelled 1 to n]
- for i=1 to n [O(n)]
    - if vertex i not yet explored [O(n), in some previous BFS]
        - BFS(G, i) [discovers precisely i's connected components]

```

Note: Finds every connected component.

Running time: $O(m + n)$, $m \rightarrow O(1)$ per edge in each BFS, $n \rightarrow O(1)$ per node

1.3 Depth First Search (DFS): The Basics

1.3.1 Overview and Example

Explore aggressively, only backtrack when necessary

Also computes a topological ordering of a directed acyclic graph

and strongly connected components of directed graphs

Running time: $O(m + n)$

1.3.2 The Code

Exercise: mimic BFS code, use a stack instead of a queue [+minor other modifications].

Recursive version:

```

DFS (graph G, start vertex s)
    - mark s as explored
    - for every edge (s, v):
        - if v unexplored
            - DFS(G, v)

```

1.3.3 Basic DFS Properties

Claim #1: at end of the algorithm, v marked as explored $\iff \exists$ path from s to v in G .

Reason: particular instantiation of generic search procedure.

Claim #2: running time is $O(n_s + m_s)$

- n_s = number of nodes reachable from s
- m_s = number of edges reachable from s

Reason: nodes at each node in connected component of s at most once, each edge at most twice.

1.4 Topological Sort

Definition: A topological ordering of a directed graph G is a labelling F of G 's nodes such that:

1. the $f(v)$'s are the set $\{1, 2, \dots, n\}$
2. $(u, v) \in G \Rightarrow f(u) < f(v)$

Motivation: sequence tasks while respecting all precedence constraints.

Note: G has directed cycle \Rightarrow no topological ordering.

Theorem: no directed cycle \Rightarrow can recompute topological ordering in $O(m + n)$ time

1.4.1 Straightforward Solution

Every directed, acyclic graph has a sink vertex, a vertex without any outgoing arcs.

Reason: if not, can keep following outgoing arcs to produce directed cycle.

To compute topological ordering:

- let v be a sink vertex of G
- set $f(v)=n$
- recurse on $G-\{v\}$

Why does it work? When v is assigned to position i , all outgoing arcs already deleted \Rightarrow all lead to later vertices in ordering.

1.4.2 Topological Sort via DFS (Slick)

DFS-Loop(graph G)

- mark all nodes unexplored
- $\text{current_label}=n$ [to keep track of ordering]
- for each vertex v in G :
 - if v not yet explored [in some previous DFS call]
 - DFS(G, v)

DFS(graph G , start vertex s)

- mark s explored
- for every edge (s, v) :
 - if v not yet explored
 - DFS(G, v)
- set $F(s)=\text{current_label}$
- $\text{current_label}--$

[16:05]

Running time: $O(m+n)$

Reason: $O(1)$ time per node, $O(1)$ time per edge

Correctness: need to show that if (u, v) is an edge, then $f(u) < f(v)$.

Case 1: u visited by DFS before $v \Rightarrow$ recursive call corresponding to v finishes before that of u (since DFS) $\Rightarrow f(v) > f(u)$.

Case 2: v visited before $u \Rightarrow v$'s recursive call finishes before u 's even starts $\Rightarrow f(v) > f(u)$

1.5 Computing Strong Components: The Algorithm

Formal Definition: The strongly connected components (SCCs) of a directed graph G are the equivalence classes of the relation $u \sim v \iff \exists \text{path } u \rightsquigarrow v \text{ and a path } u \rightsquigarrow u \text{ in } G$.

1.5.1 Why Depth-First Search

Depending on which vertex you start with, you can end up with one component - whole graph or many.

1.5.2 Kosaraju's Two-Pass Algorithm

Theorem: Can compute SCCs in $O(n)$ time

Algorithm: (given directed graph G)

1. Let $G^{reverse} = G$ with all arcs reversed
2. run DFS-Loop on G^{rev}
 - (a) goal: compute “magical ordering” of nodes
 - (b) Let $f(v)$ =”finishing time” of each $v \in V$
3. run DFS-Loop on G
 - (a) goal: discover the SCCs one-by-one
 - (b) processing nodes in decreasing order of finishing times
 - (c) [SCCs=nodes with the same “leader”]

1.5.3 DFS-Loop

```

- global variable t=0 [for finishing times in 1st pass
[#of nodes processed so far
- global variable s=NULL [for leaders in 2nd pass]
[current source vertex]
Assume nodes labeled 1 to n
- for i=n downto 1
    - if i not yet explored
        - s:=i
        DFS(G, i)

DFS(graph G, node i)
- mark i as explored [for rest of DFS loop
- set leader(i):=node s
- for each arc (i, j) in G:
    - if j not yet explored:
        - DFS(G, j)

- t++
- set f(i):=t
[f(i) i's finishing time]
```

For second pass \Rightarrow reverse orientation, change names to $f(i)$'s.

Running Time: $2 \cdot DFS = O(m + n)$

1.6 Computing Strong Components: The Analysis

1.6.1 Observation

Claim: the SCCs of a directed graph induce an acyclic “meta graph”:

mega-nodes = the SCCs C_1, \dots, C_k of G

$$\exists \text{ arc } C \rightarrow \hat{C} \iff \exists \text{ arc } (i, j) \in G \text{ with } i \in C, j \in \hat{C}$$

1.6.2 Why acyclic?

A cycle of SCCs would collapse into one.

SCC of the original graph G and its reversal G^{rev} is exactly the same.

1.6.3 Key Lemma

Lemma: [7:00] Consider two “adjacent” SCCs in G :

Let $F(v)$ = finishing times of DFS-Loop in G^{rev} .

Then: $\max_{v \in C_1} F(v) < \max_{v \in C_2} F(v)$

Corollary: maximum F-value of G must lie in a “sink SCC” [10:00]

1.6.4 Correctness Intuition

By Corollary: 2nd pass of DFS-Loop begins somewhere in a sink SCC C^* .

\Rightarrow First call to DFS discovers C^* and nothing else.

\Rightarrow rest of DFS-Loop like recursing on G with C^* deleted. [starts in a sink node of $G - C^*$]

\Rightarrow successive calls to DFS(G_{ij}) “peel off” SCCs one by one [in reverse topological order of the “meta-graph” of the SCCs]

1.6.5 Proof of Key Lemma

in G^{rev} : $C_1 \quad i \leftarrow j \quad C_2$ [still SCCs (of G^{rev})]

let v = 1st node of $C_1 \cup C_2$ reached by 1st pass of DFS-Loop (on G^{rev})

Case 1 [$v \in C_1$]: all of C_1 explored before C_2 ever reached

Reason: no paths from C_1 to C_2 (since meta-graph is acyclic)

\Rightarrow all F-values in C_1 are less than all of F-values in C_2 .

Case 2 [$v \in C_2$]: DFS(G^{rev}, v) won't finish until all of $C_1 \cup C_2$ completely explored $\Rightarrow \forall_{w \in C_1} F(v) > F(w)$.